

# CityLink Farebox (Python Version)

## Overview

CityLink operates a metro card system with billing rules that have evolved over time and were never formally documented. This project addresses the challenge of reverse-engineering these undocumented fare rules. By analyzing an anonymized tap log (containing tap times, stations, lines, and the corresponding charged amounts for a single rider), the goal is to discover the underlying fare logic. Once these rules are hypothesized, they are implemented within a flexible, object-oriented Python engine capable of accurately computing fares for new tap events.

## Problem Statement

The primary challenge is to infer a minimal and accurate set of simple fare rules that precisely reproduce the `charged ₹amount` observed in the provided historical tap log data. This inference process involves analyzing patterns in the given data against potential rule types (base fares, surcharges, discounts, transfer windows). Following rule discovery, the task is to implement these rules within a robust, object-oriented programming (OOP) engine in Python, ensuring each rule is independently testable and toggleable.

## Goal

1. **Reproduce Charges Exactly:** Through careful analysis and iterative refinement, discover the underlying fare rules that precisely match the `charged ₹amount` in the provided tap log.
2. **OOP Engine Development:** Code a small, extensible object-oriented Python engine to apply these discovered rules to any new tap records, demonstrating the inferred logic.

## Input

The application processes a series of tap log records. For demonstration purposes, these records are hardcoded as a list of `TapRecord` objects within the `main()` function. Each `TapRecord` encapsulates the essential details of a single tap event:

- **Datetime String ( `datetime_str` ):** A string representing the date and time of the tap, formatted as "MM-DD HH:MM" (e.g., "07-01 07:20"). This is parsed into a `datetime` object for time-based rule application.
- **Line ( `line` ):** A single character string identifying the metro line (e.g., "G", "R", "Y").
- **Station ( `station` ):** A two-character string representing the station code (e.g., "BD", "NC").
- **Charged Amount ( `charged_amount` ):** The actual fare that was charged for this specific tap, as provided in the anonymized log. This value is used for comparison against the `Computed Fare`.

### Example Input Record:

```
TapRecord("07-01 07:20", "G", "BD", 25)
```

# Output

For every `TapRecord` processed by the `TariffEngine`, the application prints a line to the console. This line displays the original tap record's details alongside the `Computed Fare` calculated by applying the currently active set of rules.

## Expected Output Format:

Record: {datetime\_str} | Line: {line} | Station: {station} | Charged: {charged\_amount} => Computed Fare: {computed\_fare}

## Example Output (with all rules active, values may vary based on exact rule logic and order):

```
Record: 07-01 07:20 | Line: G | Station: BD | Charged: 25 => Computed Fare: 25
Record: 07-01 08:01 | Line: G | Station: NC | Charged: 37.5 => Computed Fare: 35
Record: 07-01 08:30 | Line: R | Station: YH | Charged: 0 => Computed Fare: 35
Record: 07-01 10:01 | Line: R | Station: KL | Charged: 25 => Computed Fare: 20.0
Record: 07-01 14:36 | Line: G | Station: NC | Charged: 25 => Computed Fare: 20.0
Record: 07-01 22:15 | Line: Y | Station: BD | Charged: 20 => Computed Fare: 20.0
Record: 07-03 00:20 | Line: R | Station: NC | Charged: 16.25 => Computed Fare: 16.25
```

# Input and Output Cases

Here are a few specific scenarios demonstrating how the `CityLinkFarebox.py` application computes fares under various conditions, assuming all rules (R1-R5) are active as configured in the `main()` function.

## Assumed Active Rules:

- `BaseFareRule(active=True)` : Base fare ₹25.
- `PeakPeriodRule(active=True)` : +₹10 surcharge during 08:00-10:00 and 18:00-20:00.
- `TransferWindowRule(active=True)` : Placeholder, currently does not modify fare.
- `NightDiscountRule(active=True)` : 20% discount (x0.8) during 10:00-23:59.
- `PostMidnightRule(active=True)` : 35% discount (x0.65) during 00:00-03:59.

## Case 1: Standard Off-Peak Tap

### Input Tap Record:

`TapRecord("07-01 07:20", "G", "BD", 25)`

### Explanation of Calculation:

1. **Initial Fare:** 0
2. **BaseFareRule (R1):** Sets fare to 25 .
3. **PeakPeriodRule (R2):** Time is 07:20, which is outside peak hours (08:00-10:00, 18:00-20:00). Fare remains 25 .
4. **TransferWindowRule (R3):** Placeholder, fare remains 25 .
5. **NightDiscountRule (R4):** Time is 07:20, which is outside 10:00-23:59. Fare remains 25 .
6. **PostMidnightRule (R5):** Time is 07:20, which is outside 00:00-03:59. Fare remains 25 .

### Expected Output:

Record: 07-01 07:20 | Line: G | Station: BD | Charged: 25 => Computed Fare: 25

## Case 2: Morning Peak Period Tap

### Input Tap Record:

TapRecord("07-01 08:01", "G", "NC", 37.5)

### Explanation of Calculation:

1. **Initial Fare:** 0
2. **BaseFareRule (R1):** Sets fare to 25 .
3. **PeakPeriodRule (R2):** Time is 08:01, which is within morning peak hours (08:00-10:00). Adds 10 to fare. Fare becomes  $25 + 10 = 35$  .
4. **TransferWindowRule (R3):** Placeholder, fare remains 35 .
5. **NightDiscountRule (R4):** Time is 08:01, which is outside 10:00-23:59. Fare remains 35 .
6. **PostMidnightRule (R5):** Time is 08:01, which is outside 00:00-03:59. Fare remains 35 .

### Expected Output:

Record: 07-01 08:01 | Line: G | Station: NC | Charged: 37.5 => Computed Fare: 35

(Note: The computed fare of 35 is close to the charged 37.5, suggesting the peak surcharge might be slightly different or there's another rule/rounding involved in the original data.)

## Case 3: Night Discount Tap (Late Evening)

### Input Tap Record:

TapRecord("07-01 22:15", "Y", "BD", 20)

### Explanation of Calculation:

1. **Initial Fare:** 0
2. **BaseFareRule (R1):** Sets fare to 25 .
3. **PeakPeriodRule (R2):** Time is 22:15, which is outside peak hours. Fare remains 25 .
4. **TransferWindowRule (R3):** Placeholder, fare remains 25 .
5. **NightDiscountRule (R4):** Time is 22:15, which is within 10:00-23:59. Applies 20% discount. Fare becomes  $25 * 0.8 = 20$  .
6. **PostMidnightRule (R5):** Time is 22:15, which is outside 00:00-03:59. Fare remains 20 .

### Expected Output:

Record: 07-01 22:15 | Line: Y | Station: BD | Charged: 20 => Computed Fare: 20.0

(This case perfectly matches the charged amount, indicating the NightDiscountRule is likely correct for this scenario.)

## Case 4: Post-Midnight Discount Tap

### Input Tap Record:

TapRecord("07-03 00:20", "R", "NC", 16.25)

### Explanation of Calculation:

1. **Initial Fare:** 0
2. **BaseFareRule (R1):** Sets fare to 25 .
3. **PeakPeriodRule (R2):** Time is 00:20, which is outside peak hours. Fare remains 25 .
4. **TransferWindowRule (R3):** Placeholder, fare remains 25 .
5. **NightDiscountRule (R4):** Time is 00:20, which is outside 10:00-23:59. Fare remains 25 .
6. **PostMidnightRule (R5):** Time is 00:20, which is within 00:00-03:59. Applies 35% discount. Fare becomes  $25 * 0.65 = 16.25$  .

### Expected Output:

Record: 07-03 00:20 | Line: R | Station: NC | Charged: 16.25 => Computed Fare: 16.25

(This case also perfectly matches the charged amount, suggesting the PostMidnightRule is likely correct for this scenario.)

## Fare Rules Hypothesis (Solution Details)

Based on the problem description and typical public transport fare structures, we hypothesize the following five rules. Each rule is designed to be independently testable and contributes to the overall fare calculation.

- **R1: BaseFareRule**
  - **Description:** This rule establishes a universal base fare of ₹25 for every tap, irrespective of the line or time. This is often the starting point for any fare calculation.
  - **How it was done:** Implemented as `BaseFareRule` , its `apply` method simply returns `25` if active, effectively setting the initial fare for any tap.
  - **Testability:** By activating only this rule, one can verify that all computed fares are exactly ₹25.
- **R2: PeakPeriodRule**
  - **Description:** A surcharge of ₹10 is applied during identified peak travel periods: 08:00 to 10:00 (morning rush) and 18:00 to 20:00 (evening rush). This accounts for higher demand during these times.
  - **How it was done:** The `PeakPeriodRule` checks the `hour` component of the `tap_record.datetime` . If it falls within the defined peak hour ranges, it adds `10` to the `current_fare` .
  - **Testability:** Taps scheduled within these windows should show an additional ₹10 compared to off-peak taps (assuming `BaseFareRule` is also active).
- **R3: TransferWindowRule**
  - **Description:** Taps made within a 30-minute window of a previous tap are considered free transfers, meaning no additional fare is charged for the subsequent tap. This encourages multi-leg journeys.
  - **How it was done:** This rule is currently a **placeholder**. A complete implementation would require the `TariffEngine` to maintain state about the *previous* tap's time, or for the `apply` method to receive a list of recent taps. The current `apply` method simply returns the `current_fare` unchanged.
  - **Testability:** (Once implemented) One would simulate two taps, one within 30 minutes of the other, and expect the second tap's fare to be ₹0 or unchanged from the previous tap's fare.

- **R4: NightDiscountRule**
  - **Description:** A 20% discount is applied to the fare for rides taken during the broader night period, specifically between 10:00 (10 AM) and 23:59 (midnight). This might be to encourage off-peak travel or reflect lower demand.
  - **How it was done:** The `NightDiscountRule` checks if `tap_record.datetime.hour` is between 10 (inclusive) and 24 (exclusive). If so, it multiplies the `current_fare` by 0.8 (representing a 20% discount).
  - **Testability:** Taps within this time frame should result in a fare that is 80% of what it would be without this discount.
- **R5: PostMidnightRule**
  - **Description:** A more significant 35% discount is applied for rides taken in the very early hours, from 00:00 (midnight) to 03:59 (4 AM). This targets very low-demand periods.
  - **How it was done:** The `PostMidnightRule` checks if `tap_record.datetime.hour` is between 0 (inclusive) and 4 (exclusive). If true, it multiplies the `current_fare` by 0.65 (representing a 35% discount).
  - **Testability:** Taps in this specific early morning window should result in a fare that is 65% of the base fare (or base + surcharge, if applicable).

## Solution Approach (OOP Engine)

The solution leverages an object-oriented design in Python to create a flexible, modular, and easily maintainable fare calculation engine. This approach directly addresses the requirement for implementing rules as small classes and chaining them.

## Class Design Details

- **TapRecord Class:**
  - **Purpose:** To provide a clean, structured way to represent each line of the raw tap log data. Instead of passing around raw strings or tuples, a `TapRecord` object bundles all relevant information (datetime, line, station, original charged amount) into a single, self-contained unit.
  - **Implementation:** It includes a `datetime.strptime` call in its `__init__` method to immediately convert the `datetime_str` into a `datetime` object, making time-based calculations easier for the rules. The `__str__` method provides a user-friendly representation for output.
- **FareRule (Base Class):**
  - **Purpose:** This abstract base class (though not formally using `abc` module, it acts as one) defines the common interface that all specific fare rules must adhere to. This is crucial for implementing the "chain of responsibility" pattern.
  - **Implementation:** It has an `__init__` method that takes an `active` boolean flag, allowing each rule to be individually enabled or disabled. The `apply` method is defined to take a `tap_record` and the `current_fare` (the fare calculated by preceding rules) and is expected to return the updated fare. This design ensures polymorphism: the `TariffEngine` doesn't need to know the specific type of rule, only that it has an `apply` method.
- **Concrete FareRule Classes (e.g., BaseFareRule , PeakPeriodRule ):**
  - **Purpose:** Each of these classes inherits from `FareRule` and encapsulates the specific logic for one hypothesized fare rule. This adheres to the Single Responsibility Principle, making each rule easy to understand, test, and modify in isolation.
  - **Implementation:** Each class overrides the `apply` method from the base class. Inside `apply`, it first checks `if not self.active` to allow for toggling. Then, it implements its specific fare calculation logic (e.g., setting a base, adding a surcharge, applying a discount) based on the `tap_record`'s attributes and the `current_fare`.
- **TariffEngine Class:**
  - **Purpose:** This class acts as the central orchestrator. Its responsibility is to take a list of `FareRule` objects and apply them sequentially to a given `TapRecord` to compute the final fare. This implements the "chaining" of rules.

- **Implementation:** The `__init__` method takes a list of `FareRule` instances. The `compute_fare` method iterates through this list. For each rule, it calls its `apply` method, passing the `tap_record` and the `fare` calculated so far. The `fare` variable is updated in each step, ensuring that rules are applied in the order they appear in the list.

## Toggleable Rules (A/B Testing Hypothesis)

A core requirement was the ability to quickly test different hypotheses by enabling or disabling individual rules. This is achieved directly through the `active` boolean flag in the `FareRule` base class.

- **How it was done:** When `FareRule` instances are created in the `main()` function, they are passed `active=True` or `active=False`. Each rule's `apply` method explicitly checks this `self.active` flag at the very beginning. If `active` is `False`, the rule immediately returns the `current_fare` without applying its logic, effectively bypassing itself in the calculation chain.
- **Benefit:** This allows for rapid experimentation. For example, one could disable all discount rules to see the base fare plus peak surcharges, or enable only one discount rule at a time to isolate its effect on the charges. This is invaluable for debugging and validating the inferred rules against the historical data.

## How to Execute

### Requirements

- Python 3.6 or higher.

### Run the Application

1. **Save the Code:** Ensure the provided Python code is saved as `CityLinkFarebox.py` in your project directory.
2. **Open Terminal:** Open a command prompt or terminal window.
3. **Navigate:** Change your current directory to the folder where `CityLinkFarebox.py` is located.
4. **Execute:** Run the following command:

```
python CityLinkFarebox.py
```

The application will then process the sample tap records defined in the `main()` function and print the computed fare for each, based on the currently active rules.

## Current Status and Next Steps

The project has successfully established a robust, object-oriented framework for fare calculation based on a set of hypothesized rules.

- **Achieved:**
  - A modular OOP engine is fully implemented.
  - Hypothesized rules (R1: Base Fare, R2: Peak Period Surcharge, R4: Night Discount, R5: Post-Midnight Discount) are functional and integrated.
  - Each rule is independently toggleable, fulfilling the A/B testing requirement.
  - The system provides a clear output of computed fares for each tap record.
- **Pending Refinement:**

- **TransferWindowRule (R3):** This rule is currently a placeholder. To accurately reproduce the charges, its implementation needs to be enhanced. This would involve the `TariffEngine` (or a higher-level context) maintaining state about the *last tap's time* for the current rider, or passing a history of taps to the `apply` method of the `TransferWindowRule` so it can determine if the 30-minute window applies.
- **Exact Charge Reproduction:** While the framework is complete, the exact parameters (e.g., specific surcharge amounts, discount percentages, precise time boundaries) of the rules may need further calibration against the *entire* anonymized tap log data to achieve 100% reproduction of historical charges. This iterative process of comparing computed fares with actual charged amounts will help fine-tune the rules.
- **Future Enhancements:**
  - **External Data Loading:** Implement functionality to load tap records from an external file (e.g., CSV) instead of hardcoding them.
  - **Rule Order Optimization:** Investigate if the order of rule application significantly impacts the final fare and if an optimal order can be determined.
  - **Complex Rule Handling:** Extend the framework to handle more complex rules, such as daily caps, weekly passes, or multi-zone fares, if discovered.